

Team



Pawel Biernacki, Andreas Bleuler,
Sebastien Deldon, Claudio Gheller,
Andreas Jocksch, Douglas Potter,
Romain Teyssier

Problem trying to solve

- Astrophysics hydrodynamics for star formation and galaxy evolution based on Adaptive Mesh Refinement approach
- Speeding up hydro kernel computations - GPU as an accelerator approach

Prior Profile

- Focusing on hiding the GPU data copying behind the CPU work
- Kernel not computationally expensive enough

Run completed
Total elapsed time: 36.983540058135986

| seconds | % | STEP (rank= 1) |
|---------|-------|------------------|
| 0.142 | 0.3 | refine |
| 4.580 | 11.0 | load balance |
| 1.794 | 4.3 | courant |
| 0.250 | 0.6 | hydro - set unew |
| 11.916 | 28.5 | hydro - godunov |
| 0.209 | 0.5 | hydro - set uold |
| 0.215 | 0.5 | hydro - upload |
| 22.689 | 54.3 | flag |
| 41.795 | 100.0 | TOTAL |

STOP


Evolution and Strategy

- Let's do things in parallel! And refine the work by splitting it between group members!
- And then let's try even more things!




Evolution and Strategy

- Let's do things in parallel! And refine the work by splitting it between group members!
- And then let's try even more things!




Evolution and Strategy

- Let's do things in parallel! And refine the work by splitting it between group members!
- And then let's try even more things!




Evolution and Strategy

- Let's do things in parallel! And refine the work by splitting it between group members!
- And then let's try even more things!




Evolution and Strategy

- Let's do things in parallel! And refine the work by splitting it between group members!
- And then let's try even more things!



Evolution and Strategy

- Let's do things in parallel! And refine the work by splitting it between group members!
- And then let's try even more things!



Results and Final Profile

```
File Edit View Search Terminal Tabs Help
Terminal
Main step= 10 mcons= 0.00E+00 econs= 1.65E-14 epot= 0.00E+00 ekin= 1.25E-01
Fine step= 10 t= 2.52212E-07 dt= 4.684E-08 a= 1.000E+00 mem=26.5%
Run completed
Total elapsed time: 41.02718997001648

seconds      %      STEP (rank= 1)
0.239        0.6      refine
1.664        4.0      courant
0.330        0.6      hydro - set data
9.393        22.8     hydro - godunov
0.125        0.3      hydro - set data
0.320        0.8      hydro - upload
29.144       70.7     flag
41.220      100.0     TOTAL

Warning: ieee_inexact is signaling
FORTRAN STOP
==23690== Profiling application: bin/ramses3d sedov3d.nml
==23690== Profiling result:
No kernels were profiled.

==23690== API calls:
Time(%)   Time      Calls      Avg      Min      Max      Name
99.21% 179.08ms      32  5.5963ms  598ns  179.06ms  cudaEventCreateWithFlags
0.26% 477.88us      1  477.88us  477.88us  477.88us  cudaHostUnregister
0.23% 410.81us      1  410.81us  410.81us  410.81us  cudaHostRegister
0.14% 251.51us     83  3.0300us  190ns  106.96us  cuDeviceGetAttribute
0.05% 96.422us    48  2.0080us  333ns  6.7650us  cuPointerGetAttribute
0.03% 61.389us     3  20.463us  16.257us  27.293us  cudaStreamCreate
0.03% 47.111us    32  1.4720us  535ns  28.576us  cudaEventDestroy
0.02% 34.283us     1  34.283us  34.283us  34.283us  cuDeviceTotalMem
0.02% 29.086us     1  29.086us  29.086us  29.086us  cuDeviceGetName
0.01% 17.526us     3  5.8420us  1.8220us  13.559us  cudaStreamDestroy
0.00% 1.6430us     2    821ns   302ns   1.3410us  cuDeviceGet
0.00% 1.6040us     2    802ns   654ns   950ns     cuDeviceGetCount

hck35@daint02:/scratch/daint/hck35/mini-ramses-oct/mini-ramses-oct/mini-ramses/trun
```

```
File Edit View Search Terminal Tabs Help
Terminal
Fine step= 10 t= 2.52212E-07 dt= 4.684E-08 a= 1.000E+00 mem=26.5%
Run completed
Total elapsed time: 35.63284087181091

seconds      %      STEP (rank= 1)
0.240        0.7      refine
1.654        4.6      courant
4.017        11.2     hydro - godunov
0.320        0.9      hydro - upload
29.137       81.3     flag
35.825      100.0     TOTAL

Warning: ieee_inexact is signaling
FORTRAN STOP
==23586== Profiling application: bin/ramses3d sedov3d.nml
==23586== Profiling result:
Time(%)   Time      Calls      Avg      Min      Max      Name
42.76% 770.35ms    1560  493.81us  2.2400us  6.1190ms  [CUDA memcpy HtoD]
24.97% 449.95ms     80  5.6244ms  5.6230ms  5.6459ms  [CUDA memcpy DtoH]
6.10% 109.93ms    240  458.04us  456.01us  460.93us  godfine1_577_gpu
5.54% 99.818ms     80  1.2477ms  1.2392ms  1.2582ms  trace3d_525_gpu
2.27% 40.887ms     80  511.09us  510.69us  511.33us  uslope_989_gpu
1.02% 18.343ms     80  229.29us  227.49us  232.00us  ctoprim_872_gpu
1.01% 18.256ms     80  228.20us  227.97us  228.48us  unsplit_125_gpu
1.01% 18.205ms     80  227.56us  227.17us  228.55us  unsplit_153_gpu
0.99% 17.923ms     80  224.04us  223.84us  224.29us  unsplit_98_gpu
0.59% 10.695ms    240  44.561us  41.696us  50.432us  godfine1_532_gpu
0.43% 7.7025ms     80  96.280us  96.097us  96.417us  unsplit_135_gpu
0.43% 7.6933ms     80  96.165us  95.841us  96.833us  unsplit_164_gpu
0.39% 7.0219ms     80  87.773us  87.553us  87.969us  unsplit_108_gpu

==23586== API calls:
Time(%)   Time      Calls      Avg      Min      Max      Name
```

Results and Final Profile

==11336== Profiling application: /scratch/daint/ajocksch/Hackathon/ramses3d sedov3d.nml

==11336== Profiling result:

| Time(%) | Time | Calls | Avg | Min | Max | Name |
|---------|----------|-------|----------|----------|----------|--------------------|
| 43.60% | 576.72ms | 320 | 1.8022ms | 329.16us | 2.5166ms | [CUDA memcpy DtoH] |
| 20.71% | 273.94ms | 750 | 365.25us | 2.2400us | 2.0191ms | [CUDA memcpy HtoD] |
| 17.13% | 226.64ms | 240 | 944.34us | 919.53us | 982.99us | cmpflxm_757_gpu |
| 7.51% | 99.327ms | 80 | 1.2416ms | 1.2317ms | 1.2541ms | trace3d_519_gpu |
| 3.13% | 41.416ms | 80 | 517.70us | 517.26us | 517.96us | uslope_985_gpu |
| 1.46% | 19.277ms | 80 | 240.96us | 230.66us | 248.45us | ctoprim_868_gpu |
| 1.35% | 17.894ms | 80 | 223.67us | 223.24us | 223.97us | unsplit_151_gpu |
| 1.34% | 17.751ms | 80 | 221.89us | 221.64us | 222.18us | unsplit_96_gpu |
| 1.34% | 17.723ms | 80 | 221.53us | 221.28us | 221.86us | unsplit_123_gpu |
| 0.77% | 10.200ms | 240 | 42.499us | 40.384us | 46.753us | godfine1_2_672_gpu |
| 0.57% | 7.5363ms | 80 | 94.203us | 93.793us | 94.402us | unsplit_161_gpu |
| 0.57% | 7.4813ms | 80 | 93.516us | 93.089us | 93.730us | unsplit_133_gpu |
| 0.52% | 6.9374ms | 80 | 86.717us | 86.401us | 87.010us | unsplit_106_gpu |

Time of memcpy hidden behind the CPU!!!

What problems you encountered

- Problems with legacy app structure
 - too many nested calls and rigid call structure
- Issues with algorithm
 - flexible data format, but problem not computationally expensive enough
- Random access to memory due to the data structure
- [TODO] Communication for boundaries from within the kernel - needs rethinking

Wishlist

- magic compiler flag: “-faster” ;)
- nvprof is great, but depends on CUDA version
- cutting-edge tools available at the moment of starting the hackathon (CUDA7.5 and pgprof relation)
- interchangeability of compilers
- BUT: a sneak peek at a new PGI compiler