



Rocket

Mauro Bianco, Carlos Osuna, Paolo
Crosetto, Stefan Moosbrugger, Vassilis
Karakasis, Jacob Progsch

Problem trying to solve

- GridTools Libraries implementation:
 - For solving PDEs in weather and climate.
 - For “generic” numerical algorithms on structured grids

Prior Profile

- Focus on an horizontal and a vertical stencils used in weather codes:
 - Fourth order horizontal diffusion operator (finite differences)
 - Vertical advection (implicit tridiagonal solve)

Evolution and Strategy

- Baseline:
 - smem for intermediate results
 - Global load through texture
- Experiments
 - K-parallel (vertical)=>(5% to 60%)
 - Prefetching in smem=>(no)
 - Prefetching in regs=>(no)
 - Cache in registers=>(30%)
 - Using double2=>(0% to 21%)
 - Const mem for strides=>(17%)
 - Warp assignment=>(no)
 - Cache bank size=>(13%)

Results and Final Profile

- Accomplished overall speedup of more than 50% on a specific case
- How easy and how general will be the implementation inside GridTools?
- How will these result evolve on newer architectures?

What problems you encountered

- We were expecting the compiler to do some of these things (register cache, efficient unrolling)
- The baseline concerning the warp assignment was already optimal
- Instruction Level Profiler is already available (CUDA 7.5) but not with Kepler

Wishlist

- CUDA compiler:
 - Why `__device__`, `__host__` on user functions?
 - Support for more STL header functions
 - Support recent standards (recently quite some improvements)
 - Public roadmap: when will feature become available?
 - Kill warnings for calls to “`__host__ __device__`”
 - Faster compilation?